Image Recognition for Medical Diagnosis

Submitted in partial fulfillment of the requirements for the degree of

Bachelor of Technology

in

Programme

by

VAISHNAV NALLADATH 20BCE2535

PATTETI REEHA MILAN 20BCE2707

NITESH LOHIYA 20BCE2747

Under the guidance of

Dr. Lijo V.P

School of Computer Science and Engineering,

VIT, Vellore.

May, 2024

<u>DECLARATION</u>

I hereby declare that the thesis entitled "Image Recognition for Medical Diagnosis" submitted by me, for the award of the degree of *Bachelor of Technology in Programme* to VIT is a record of bonafide work carried out by me under the supervision of Dr. Lijo V.P.

I further declare that the work reported in this thesis has not been submitted and will not be submitted, either in part or in full, for the award of any other degree or diploma in this institute or any other institute or university.

Place : Vellore

Date :

Signature of the Candidate

CERTIFICATE

This is to certify that the thesis entitled "Image Recognition for Medical Diagnosis" submitted by **Vaishnav Nalladath (20BCE2535), Patteti Reeha Milan (20BCE2707), Nitesh Lohiya (20BCE2747), School of Computer Science and Engineering**, VIT, for the award of the degree of *Bachelor of Technology in Programme*, is a record of bonafide work carried out by them under my supervision during the period, 01. 12. 2023 to 30.04.2024, as per the VIT code of academic and research ethics.

The contents of this report have not been submitted and will not be submitted either in part or in full, for the award of any other degree or diploma in this institute or any other institute or university. The thesis fulfills the requirements and regulations of the University and in my opinion meets the necessary standards for submission.

Place : Vellore

Date :

Signature of the Guide

Internal Examiner

External Examiner

Head of the

Department

Programme

ACKNOWLEDGEMENTS

I wish to express my sincere thanks to my guide Dr. Lijo V.P, Assistant Professor Sr. Grade 1, Vellore Institute of Technology, Vellore for allowing me to partake in this project review as part of my CAPSTONE project which has proven to be helpful and is certain to be useful for my future career. I also wish to extend my thanks to the school office (SCOPE), HOD Dr. Umadevi K S and the dean Dr. Ramesh Babu K of VIT Vellore for allowing for the facilitation of this project for my benefit.

Student Name

Executive Summary

The "Image Recognition for Medical Diagnosis" project represents a pioneering effort aimed at reshaping healthcare paradigms through the fusion of cutting-edge artificial intelligence (AI) and image recognition technologies. By leveraging deep learning algorithms and an extensive repository of medical imagery encompassing X-rays, MRIs, and CT scans, our objective is to elevate the precision, velocity, and accessibility of medical diagnostics. Through automation of disease detection and categorization, we aspire to diminish reliance on subjective human interpretation and mitigate diagnostic errors. Our primary focus lies in crafting an intricate AI framework trained on meticulously annotated medical images, designed to discern subtle patterns and anomalies, thereby refining diagnostic accuracy across diverse medical domains. Furthermore, we confront ethical, privacy, and integration hurdles to ensure the seamless, secure, and patient-centric implementation of this transformative technology within healthcare ecosystems. This initiative heralds a quantum leap towards the digital metamorphosis of healthcare, heralding a future marked by augmented diagnostic prowess, proactive disease interception, and ultimately, superior patient outcomes, emblematic of the symbiotic alliance between technology and medicine

	CONTENTS	Page No.
	Acknowledgement	**i.**
	Executive Summary	**ii.**
	Table of Contents	**iii.**
	List of Figures and Tables	**v.**
1	**INTRODUCTION**	**1**
	1.1 Objectives	**1**
	1.2 Motivation	**1**
	1.3 Background	**1**
2	**PROJECT DESCRIPTION AND GOALS**	**2**
	2.1 Survey on Existing System	**2**
	2.2 Research Gap	**6**
	2.3 Problem Statement	**6**
3	**TECHNICAL SPECIFICATION**	**7**
	3.1 Requirements 3.1.1 Functional 3.1.2 Non-Functional	**7** **7** **8**
	3.2 Feasibility Study 3.2.1 Technical Feasibility 3.2.2 Economic Feasibility 3.2.3 Social Feasibility	**10** **10** **10** **11**
	3.3 System Specification	**11**

	3.3.1 Hardware Specification 3.3.2 Software Specification 3.3.3 Standards and Policies	**11** **12** **12**
4	**DESIGN APPROACH AND DETAILS**	**14**
	4.1 System Architecture	**14**
	4.2 Design 4.2.1 Data Flow Diagram 4.2.2 Use Case Diagram 4.2.3 Class Diagram 4.2.4 Sequence Diagram	**15** **15** **16** **17** **17**
	4.3 Constraints, Alternatives and Tradeoffs	**18**
5	**SCHEDULE, TASKS AND MILESTONES**	**19**
	5.1 Gantt Chart	**19**
	5.2 Module Description 5.2.1 Data Ingestion Module 5.2.2 Data Storage and Management 5.2.3 Image Processing and Feature Extraction 5.2.4 AI and Machine Learning Engine 5.2.5 Decision Support and Reporting Module	**20** **20** **20** **20** **20** **20**
6	**PROJECT DEMONSTRATION**	**21**
7	**COST ANALYSIS / RESULT & DISCUSSION**	**21**
8	**SUMMARY**	**23**
9	**REFERENCES**	**23**
	APPENDIX A – SAMPLE CODE	**26**

List of Figures

Figure No.	Title	Page No.
4.1	System Architecture	14
4.2.1	Data Flow Diagram	15
4.2.2	Use Case Diagram	16
4.2.3	Class Diagram	17
4.2.4	Sequence Diagram	17

List of Tables

Table No.	Title	Page No.
2.1	Literature Survey	2

1. INTRODUCTION

1.1 OBJECTIVE

The objective of the "Image Recognition for Medical Diagnosis" project is to develop and implement an AI-based system that accurately and efficiently analyzes medical images, such as X-rays, MRI, and CT scans, to assist healthcare professionals in diagnosing various diseases and conditions, thereby improving patient care outcomes and accessibility to expert-level diagnostics across diverse populations.

1.2 MOTIVATION

The motivation behind an image recognition project for medical diagnosis stems from the critical need to enhance the accuracy, efficiency, and accessibility of diagnostic processes in healthcare. With the increasing complexity and volume of medical data, particularly imaging data such as X-rays, CT scans, and MRIs, traditional diagnostic methods can be time-consuming and sometimes prone to human error. By integrating advanced image recognition technologies, such as deep learning and artificial intelligence, into the diagnostic process, we can provide clinicians with powerful tools that support and enhance their decision-making capabilities. This technology can help in identifying subtle patterns in images that might be overlooked by human eyes, enabling earlier and more accurate detection of diseases such as cancers, fractures, and neurological disorders. Moreover, automating routine analysis can significantly reduce diagnostic times, allowing healthcare providers to focus more on patient care and treatment planning. Ultimately, this leads to improved patient outcomes, more efficient use of medical resources, and the democratization of high-quality healthcare services, making them more accessible to underserved regions and populations.

1.3 BACKGROUND

The motivation and background of an image recognition project for medical diagnosis merge advancements in medical imaging technology with computational breakthroughs. Traditionally, medical image interpretation relied heavily on skilled radiologists, but the rise of digital imaging significantly increased both the volume and quality of available data, emphasizing the need for more sophisticated interpretation tools. The advent of deep learning, particularly Convolutional Neural Networks (CNNs), has revolutionized computer vision, providing the framework for automating and enhancing diagnostic accuracy and speed. This technology not only supports radiologists in detecting diseases more accurately and efficiently but also transforms medical diagnostics by facilitating earlier disease detection and tailored treatment strategies, ultimately optimizing healthcare management and improving patient outcomes.

2. PROJECT DESCRIPTION AND GOALS

2.1 SURVEY ON EXISTING SYSTEMS

S.NO	TITLE	AUTHOR	JOURNAL	SUMMARY	CITATION
1.	Computational Traditional Chinese Medicine diagnosis	Qi Zhang, Jianhang Zhou, Bob Zhang	Journal of Biomedical Informatics	This paper provides a comprehensive overview of the research conducted in the field of Computational Traditional Chinese Medicine (TCM) diagnosis. They discuss the challenges and opportunities in integrating modern computational approaches with traditional TCM diagnosis methods. The survey encompasses a range of topics such as data mining, machine learning, pattern recognition, and expert systems as they relate to TCM diagnosis. Through their analysis, the authors aim to provide insights into the current state of research and potential directions for future studies in this interdisciplinary field.	Zhang, Q., Zhou, J., & Zhang, B. (Year). Computational Traditional Chinese Medicine diagnosis: A literature survey. Journal Name, Volume (Issue), Page Range. https://doi.org/10.1007/s12065-020-00540-3
2.	Computer-aided medical diagnosis	W. Rogers, B. Ryack, G. Moeller	Journal of Medical Systems	The paper presents a comprehensive literature review on computer-aided medical diagnosis. It likely covers various aspects of using computational methods and technology to assist medical professionals in diagnosing diseases or conditions. The review may include discussions on the development, applications, challenges, and prospects of computer-aided diagnosis systems in the medical field.	Rogers, W., Ryack, B., & Moeller, G. (Year of publication not provided). Computer-aided medical diagnosis: Literature review. Journal of Medical Systems.
3.	Deep Learning	Xiaoyan	Cancers	The paper likely explores the application of	Jiang, X., Hu, Z.,

	for Medical Image-Based Cancer Diagnosis	Jiang, Zuojin Hu, Shuihua Wang, Yudong Zhang		deep learning techniques in diagnosing cancer based on medical imaging data. This paper might present case studies or experiments demonstrating the effectiveness of deep learning approaches in accurately diagnosing different types of cancers from medical images like X-rays, MRIs, CT scans, or histopathological images.	Wang, S., & Zhang, Y. (2023). Deep Learning for Medical Image-Based Cancer Diagnosis. Cancers, 15(14), 3608. https://doi.org/10.3390/cancers15143608
4.	Deep Learning in Skin Disease Image Recognition: A Review	Ling-Fang Li, Xu Wang, Wei-Jian Hu, Neal N. Xiong, Yong-Xing Du		provides a comprehensive overview and analysis of the application of deep learning techniques in the field of skin disease image recognition. They discuss the challenges and opportunities in this domain, including issues related to dataset size, model generalization, and interpretability. Additionally, the paper highlights the potential impact of deep learning approaches on improving diagnosis and treatment decision-making in dermatology.	Li, L.-F., Wang, X., Hu, W.-J., Xiong, N. N., & Du, Y.-X. (2020). Deep Learning in Skin Disease Image Recognition: A Review. Journal of Medical Systems, Volume 8
5.	Efficient and privacy-preserving medical research support platform against COVID-19: A blockchain-based approach	K. Yu, L. Tan, X. Shang, J. Huang, G. Srivastava and P. Chatterjee	EEE Consum. Electron	This paper provides a comprehensive overview of the research conducted in the field of Computational Traditional Chinese Medicine (TCM) diagnosis. They discuss the challenges and opportunities in integrating modern computational approaches with traditional TCM diagnosis methods. The survey encompasses a range of topics such as data mining, machine learning, pattern recognition, and expert systems as they relate to TCM diagnosis. Through their analysis, the authors aim to provide insights into the current state of research and potential directions for future studies in this interdisciplinary field.	Zhang, Q., Zhou, J., & Zhang, B. (Year). Computational Traditional Chinese Medicine diagnosis: A literature survey. Journal Name, Volume (Issue), Page Range. https://doi.org/10.1007/s12065-020-00540-3
6.	Secure IoT-based modern healthcare system with fault-tolerant decision-making process	P. Gope, Y. Gheraibia, S. Kabir and B. Sikdar	IEEE J. Biomed. Health Informat.	This paper presents a novel framework for a secure Internet of Things (IoT)-based healthcare system designed to enhance modern healthcare delivery. This system emphasizes the importance of security and fault tolerance in the decision-making process to ensure reliable and uninterrupted healthcare services. It proposes a mechanism that integrates advanced cryptographic techniques to safeguard patient data and ensure privacy while employing a fault-tolerant decision-making process that can operate effectively even in the event of system failures or errors.	P. Gope, Y. Gheraibia, S. Kabir and B. Sikdar, "A secure IoT-based modern healthcare system with fault-tolerant decision-making process", IEEE J. Biomed. Health Informat., vol. 25, no. 3, pp. 862-873, Mar. 2021.

7.	Convolutional sparse representation and local density peak clustering for medical image fusion	L. Wang, C. Shi, S. Lin, P. Qin and Y. Wang	Int. J. Pattern Recognit. Artif. Intell	The methodology combines convolutional sparse representation (CSR) with local density peak clustering (LDPC) to enhance the quality and effectiveness of fused medical images. Convolutional sparse representation is utilized to capture the intrinsic features and details of medical images, ensuring that critical information is preserved during the fusion process. The local density peak clustering method is applied to improve the selection of salient features by identifying and emphasizing areas of high informational density, which are crucial for accurate medical diagnosis.	L. Wang, C. Shi, S. Lin, P. Qin and Y. Wang, "Convolutional sparse representation and local density peak clustering for medical image fusion", Int. J. Pattern Recognit. Artif. Intell., vol. 34, no. 7, 2020.
8.	A survey on active learning and human-in-the-loop deep learning for medical image analysis	S. Budd, E. C. Robinson, and B. Kainz	Med. Image Anal	The paper systematically reviews various techniques and methodologies in this area, discussing their advantages, limitations, and applicability to different medical imaging tasks such as diagnosis, segmentation, and detection. It also explores how these approaches can reduce the annotation burden on medical experts, improve learning efficiency, and potentially lead to more accurate and robust deep learning models in medical image analysis. It discusses the challenges and future directions in this field, emphasizing the need for more interactive and intelligent systems that can effectively integrate human feedback and adaptively improve over time.	S. Budd, E. C. Robinson, and B. Kainz, "A survey on active learning and human-in-the-loop deep learning for medical image analysis", Med. Image Anal., vol. 71, 2021.
9.	Spherical-patches extraction for deep-learning-based critical points detection in 3D neuron microscopy images	W. Chen	IEEE Trans. Med. Imag	The paper discusses a novel approach for detecting critical points in 3D neuron microscopy images using deep learning techniques. Critical points in the context of neuron microscopy might refer to key features or structures within the neuronal architecture, such as branching points, endpoints, or junctions, which are crucial for understanding the complex connectivity and functionality of the brain. The method probably involves extracting spherical patches around potential critical points in the 3D images. These spherical patches are then used as input for a deep learning model designed to accurately classify and identify the critical points.	W. Chen et al., "Spherical-patches extraction for deep-learning-based critical points detection in 3D neuron microscopy images", IEEE Trans. Med. Imag., vol. 40, no. 2, pp. 527-538, Feb. 2021.
10.	A Method of the Breast Cancer Image Diagnosis Using Artificial Intelligence	Daewon Kwak, Jiwoo Choi, Sungjin	THE JOURNAL OF KOREAN INSTITUTE	The recent advance of image recognition technology comes from the accumulation of numerous data and deepening of neural networks. However, training these various data on a deep neural network causes	Kwak, D. and Choi, J., 2023. A Method of Breast Cancer Image Diagnosis Using

	Medical Images Recognition Technology Network	Lee	OF COMMUNICATIONS AND INFORMATION SCIENCES	various problems. Overfitting is caused by a small amount of data, class imbalance resulting from the difference in the amount of data between classes, and multi-class training problems. This paper found and analyzed these problems occurring in such small data sets, and suggested solutions and analyzed the performance through experiments. For these goals, we compared open small data sets and the differences between them and selected the training techniques that perform well for each dataset.	Artificial Intelligence Medical Images Recognition Technology Network. The Journal of Korean Institute of Communications and Information Sciences, 48(2), pp.216-226.
11.	Machine learning and deep learning approach for medical image analysis: diagnosis to detection	Meghavi Rana, Megha Bhushan	School of Computing, DIT University, Dehradun, India	Computer-aided detection in medicine, leveraging Deep Learning (DL) and Machine Learning (ML), has seen remarkable growth. Medical images yield crucial diagnostic information, vital for early disease detection and reduced mortality rates from cancer and tumors. ML faces limitations with extensive data, whereas DL handles data efficiently. DL, an advanced ML technique, learns machine behaviors. Using multilayered neural networks, DL extracts comprehensive dataset insights. This literature review covers 40 studies from Jan 2014–2022, exploring ML and DL applications in disease detection and classification. It assesses diverse approaches, imaging modalities, evaluation tools, and techniques. MRI datasets enable comparative analyses between ML classifiers and DL models, aiding healthcare decisions.	Rana, M., Bhushan, M. Machine learning and deep learning approach for medical image analysis: diagnosis to detection. Multimed Tools Appl 82, 26731–26769 (2023). https://doi.org/10.1007/s11042-022-14305-w
12.	Hybrid Intelligence-Driven Medical Image Recognition for Remote Patient Diagnosis in Internet of Medical Things	Zhiwei Guo, Yu Shen; Shaohua Wan, Wen-Long Shang, Keping Yu	IEEE Journal of Biomedical and Health Informatics	This paper proposes a hybrid intelligence-driven medical image recognition framework in IoMT for smart cities. It combines deep learning with conventional machine learning (CML) to extract deep and abstract features for initial images. The convolution neural network extracts these features, while CML-based techniques reduce dimensions for extracted features and construct a strong classifier for recognition results. A real dataset about pathologic myopia is used to establish a simulative scenario to assess the proposed recognition framework. Results show that the proposed framework improves recognition accuracy by two to three percent. This approach is aimed at improving remote patient diagnosis in IoMT.	Z. Guo, Y. Shen, S. Wan, W. -L. Shang and K. Yu, "Hybrid Intelligence-Driven Medical Image Recognition for Remote Patient Diagnosis in Internet of Medical Things," in IEEE Journal of Biomedical and Health Informatics, vol. 26, no. 12, pp. 5817-5828, Dec. 2022, doi: 10.1109/JBHI.202

					1.3139541.

2.1 GAPS IDENTIFIED

1. Limited Diversity in Datasets: Many studies note a lack of diversity in medical image datasets, leading to biases and reduced generalizability of AI models across different demographics and medical conditions.
2. Interpretability and Explainability: While AI models show promising accuracy, their lack of interpretability raises concerns about the transparency and trustworthiness of their diagnoses, especially in critical medical decisions.
3. Integration Challenges: Integrating AI-driven image recognition systems into existing healthcare infrastructure presents technical hurdles such as compatibility issues, data interoperability, and workflow disruptions.
4. Data Privacy and Security: There is a need for robust data privacy measures to protect sensitive patient information, especially when dealing with large-scale medical image datasets and AI algorithms.
5. Validation and Clinical Adoption: The gap between research validation and real-world clinical adoption remains significant, requiring more studies on the efficacy, reliability, and cost-effectiveness of AI-based diagnostic tools.

2.2 PROBLEM STATEMENT

The primary challenge this project aims to address is the need for a more accurate, efficient, and consistent analysis of medical images to aid in diagnosis. Despite the advancements in medical imaging technologies, the reliance on human interpretation can introduce subjectivity, potentially leading to diagnostic errors. These errors can have significant implications for patient care, including delayed treatment, unnecessary interventions, and increased patient anxiety. Furthermore, the growing demand for medical imaging services exacerbates the workload on radiologists, resulting in longer waiting times for patients and increasing the likelihood of burnout among healthcare professionals.

3. TECHNICAL SPECIFICATIONS

3.1 REQUIREMENTS

3.1.1 Functional

1. Image Upload and Management

The system must allow users (healthcare professionals, radiologists, etc.) to securely upload medical images in various formats (e.g., X-ray, MRI, CT scans, dermatological photographs).
The system should provide the capability to categorize, tag, and store images based on patient ID, date, type of image, and other relevant metadata for easy retrieval and reference.

2. Pre-processing of Images

Automatic image pre-processing to standardize images for analysis, including resizing, normalization, and enhancement to improve image quality and recognition accuracy.

3. Image Analysis and Recognition

Implementation of deep learning algorithms or other machine learning models to analyze images and identify specific features or patterns indicative of medical conditions.
The system should support the recognition of a broad spectrum of diseases or medical conditions, as specified in the project scope.

4. Diagnostic Support

Provide diagnostic suggestions based on image analysis, including potential conditions identified and a confidence score or probability estimate for each suggestion.
Include references to similar cases or relevant medical literature to support diagnostic suggestions.

5. Security and Privacy

Ensure all patient data, including images and diagnostic results, are stored and transmitted securely, adhering to HIPAA or equivalent privacy standards applicable in the region of deployment.
Implement user authentication and authorization to restrict access to patient data and diagnostic tools to authorized personnel only.

6. User Interface

Provide a user-friendly interface for healthcare professionals to upload images, view analyses, and receive diagnostic suggestions.

The interface should be accessible across multiple devices, including desktops, tablets, and smartphones, to accommodate different use cases.

7. Reporting and Exporting

Capability to generate reports summarizing the diagnostic findings, which can be printed or exported in various formats for sharing with other healthcare providers or for inclusion in patient records.

3.1.2 Non-Functional

1. Performance

The system should process and analyze images quickly, with response times that meet clinical needs, allowing for timely diagnosis.

It must handle simultaneous requests from multiple users without significant degradation in performance.

2. Scalability

The system should be scalable, capable of handling an increasing number of images and users as the demand grows.

It should support scaling both vertically (adding more power to existing machines) and horizontally (adding more machines).

3. Reliability and Availability

High availability is crucial; the system should be operational and accessible 24/7, considering the critical nature of medical diagnosis.

It must be reliable, with minimal downtime and the ability to recover quickly from failures.

4. Security

Implement stringent security measures to protect sensitive patient data against unauthorized access, breaches, and leaks.

Use encryption for data at rest and in transit and adhere to industry standards and regulations like HIPAA for healthcare data.

5. Privacy

Ensure patient confidentiality by anonymizing patient data used within the system for diagnosis or training machine learning models.

Comply with privacy laws and regulations applicable to the jurisdictions in which the system operates, such as GDPR or HIPAA.

6. Usability

The user interface (UI) should be intuitive and user-friendly for medical professionals with varying levels of technical expertise.

Provide adequate documentation and training materials to support users in effectively utilizing the system.

7. Maintainability

Code and system architecture should be designed for easy maintenance and updates, allowing for new features, improvements, and bug fixes to be deployed with minimal disruption.

Provide logs and monitoring tools for diagnosing issues and optimizing performance.

8. Data Integrity

Implement measures to ensure the accuracy and consistency of patient data and diagnostic information throughout its lifecycle in the system.

9. Regulatory Compliance

The system must comply with all relevant regulatory standards for medical software and devices, including certification where applicable.

10. Environmental Considerations

For hardware-dependent setups, consider energy efficiency and environmental impact, aiming for solutions that minimize carbon footprint.

3.2 FEASIBILITY STUDY

3.2.1 Technical Feasibility

Availability and Maturity of AI Models: The project's core relies on advanced image recognition algorithms, typically powered by deep learning and artificial intelligence (AI). The technical feasibility is high in this aspect because there has been significant progress in AI technologies, particularly convolutional neural networks (CNNs) and other machine learning models tailored for image analysis.
Compliance and Security Protocols: Ensuring the system's compatibility with healthcare compliance standards (e.g., HIPAA in the United States, GDPR in Europe for patient data privacy) is non-negotiable. The technical feasibility is contingent upon implementing robust encryption, user authentication, and data handling protocols that meet these stringent requirements.

Scalability to Meet Future Demands: As the adoption of the system grows, it must scale accordingly to accommodate an increasing number of users and data without degradation in performance. The feasibility study must consider future-proofing the architecture to handle growth in data volume, user base, and computational demand. This might involve adopting scalable cloud services, microservices architecture, or other technologies that allow for flexible system expansion.

3.2.2 Economic Feasibility

Initial Development Costs: These include expenses related to research and development, purchasing, or accessing relevant datasets for training AI models, hardware and software acquisition, and hiring a skilled team of data scientists, developers, and domain experts (e.g., radiologists for annotating images). A comprehensive cost estimation must consider all these aspects to ensure the project is financially viable.

Cost Savings for End Users: The system can provide economic value by reducing the time doctors and radiologists spend analyzing images, lowering the rates of diagnostic errors, and thereby potentially decreasing the costs associated with misdiagnosis or delayed treatment. These indirect benefits can be significant selling points and contribute to the project's economic feasibility by demonstrating tangible ROI (Return on Investment) for healthcare providers.

Assessment of Market Demand: The economic feasibility must consider the demand for such a system among healthcare providers. This involves evaluating the prevalence of conditions that the system aims to diagnose, the volume of imaging studies conducted, and the willingness of healthcare institutions to invest in advanced diagnostic technologies. Market research can provide insights into current needs and future trends in medical imaging and diagnostics.

3.2.3 Social Feasibility

Building Trust with Healthcare Professionals: For a new technology to be successfully integrated into medical practice, it must gain the trust of doctors, radiologists, and other healthcare providers. This involves demonstrating not only the system's accuracy and reliability but also how it can enhance their work rather than replace it. Engaging with these professionals early in the development process, including them in testing and feedback loops, and transparently sharing performance data can help build this trust.

Addressing Ethical Concerns: The ethical implications of using AI in medical diagnosis, such as potential biases in AI algorithms or the consequences of erroneous diagnoses, are significant. Ensuring the system is developed with ethical guidelines in mind, including fairness, accountability, and transparency, is essential for social feasibility. This also involves implementing rigorous validation and testing protocols to identify and mitigate biases that may arise from the training data or algorithmic preferences, ensuring equitable treatment for all patients.

Ensuring Broad Access: The social feasibility of the project also hinges on its accessibility to a wide range of healthcare settings, including under-resourced clinics and hospitals in rural or economically disadvantaged areas. This means the system should be affordable and require minimal infrastructure to implement, ensuring that the benefits of advanced diagnostic technologies are not limited to well-funded urban centers.

3.3 System Specification

3.3.1 Hardware Specification

There is not much required in terms of hardware except for a decent laptop with x64 architecture, minimum 8 GB of RAM and an above average Processor installed.

3.3.2 Software Specification

To run Python with PyTorch and CUDA, the following requirements are needed:

Python

Python 3.6 or later installed on the system. The latest version of Python can be downloaded from the official website: https://www.python.org/downloads/

PyTorch

This can be installed using pip by running the following command in the terminal: pip install torch torchvision

CUDA

To use PyTorch with CUDA, a compatible NVIDIA GPU is needed and the CUDA toolkit should be installed on the system. This can be downloaded from the official NVIDIA website: https://developer.nvidia.com/cuda-downloads

cuDNN

The cuDNN library, which is a CUDA-accelerated library for deep neural networks is also required. This can be downloaded from the NVIDIA website: https://developer.nvidia.com/cudnn

3.3.3 Standards and Policies

1. Data Privacy and Security

HIPAA (Health Insurance Portability and Accountability Act): In the United States, HIPAA sets the standard for protecting sensitive patient data. Any entity dealing with protected health information (PHI) must ensure that all the required physical, network, and process security measures are in place and followed.
GDPR (General Data Protection Regulation): For projects operating in or serving users in the European Union, GDPR imposes strict rules on data protection and privacy. This includes obtaining explicit consent from individuals regarding the use of their personal data, the right to access their data, and the 'right to be forgotten'.

2. Medical Device Regulations

FDA Regulations for Software as a Medical Device (SaMD): In the United States, the FDA

classifies certain software, including some AI-driven diagnostic tools, as Medical Devices, which must comply with specific regulatory requirements for approval.
EU Medical Device Regulation (MDR): Similar to the FDA, the European Union has its set of regulations for medical devices, including software. Compliance with these regulations is necessary for market approval in the EU.

3. Standards for Clinical Validity and Reliability

ISO 13485: This is a standard for a Quality Management System (QMS) specifically designed for the design and manufacture of medical devices, ensuring that products consistently meet customer and regulatory requirements.
IEC 62304: A standard for the life cycle management of medical device software, focusing on the development, maintenance, and risk management of medical software.

4. Interoperability Standards

HL7 (Health Level Seven International) and FHIR (Fast Healthcare Interoperability Resources): These standards are crucial for the integration of new medical software with existing healthcare information systems, ensuring seamless communication and data exchange between disparate system

4. DESIGN APPROACH AND DETAILS

4.1 SYSTEM ARCHITECTURE

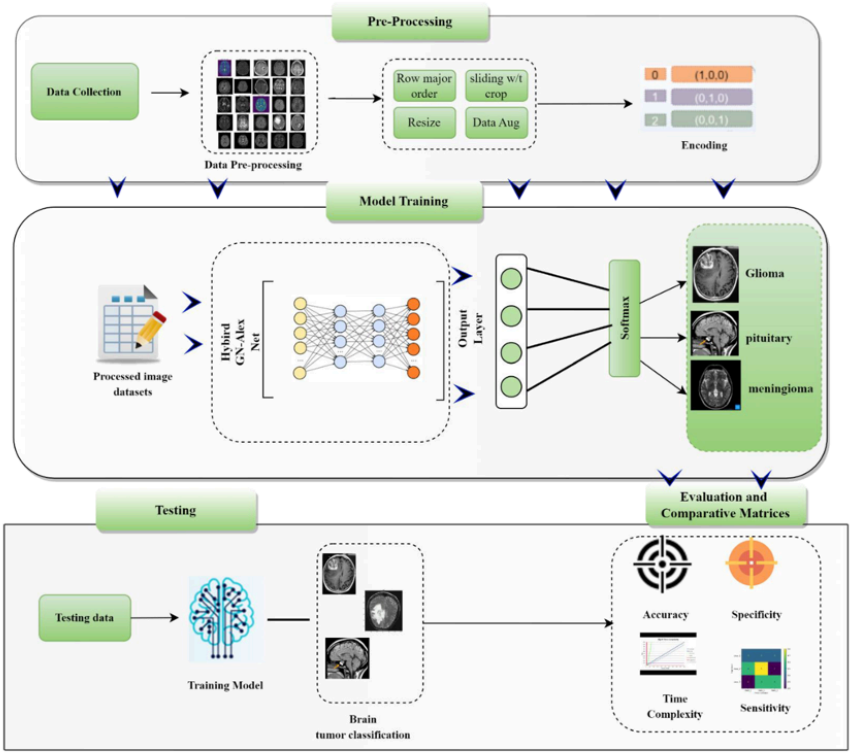

Figure 1 : System Architecture [13]

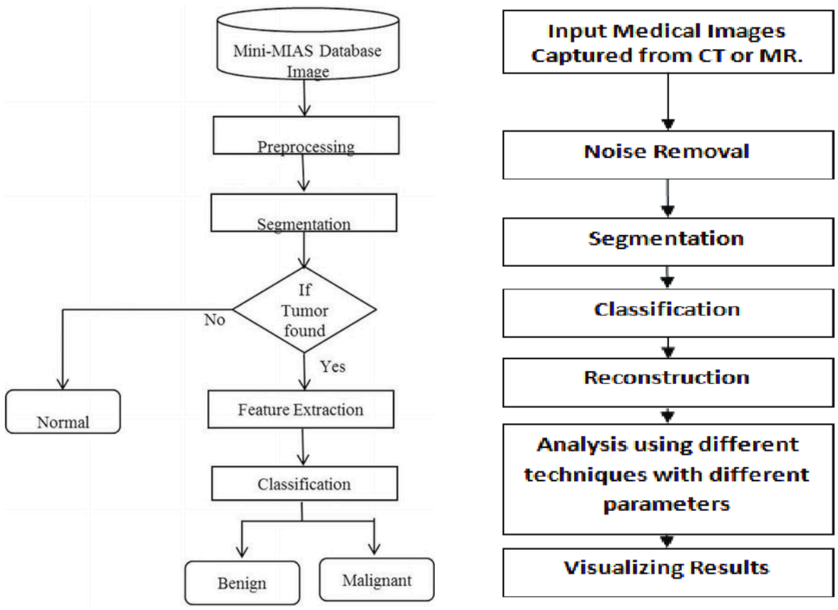

Figure 2: System Architecture (Working Principle)

4.2 DESIGN

4.2.1 DATA FLOW DIAGRAM

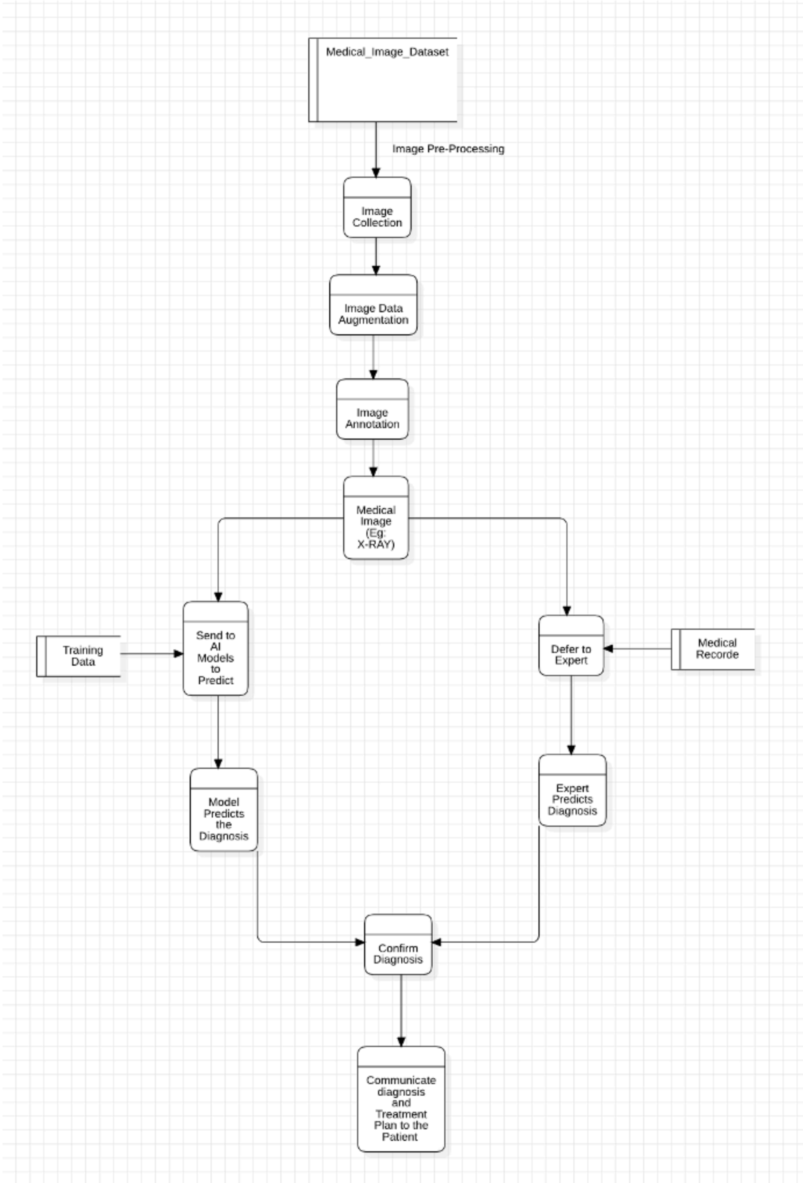

Figure 3 : Data Flow Diagram

4.2.2 USE CASE DIAGRAM

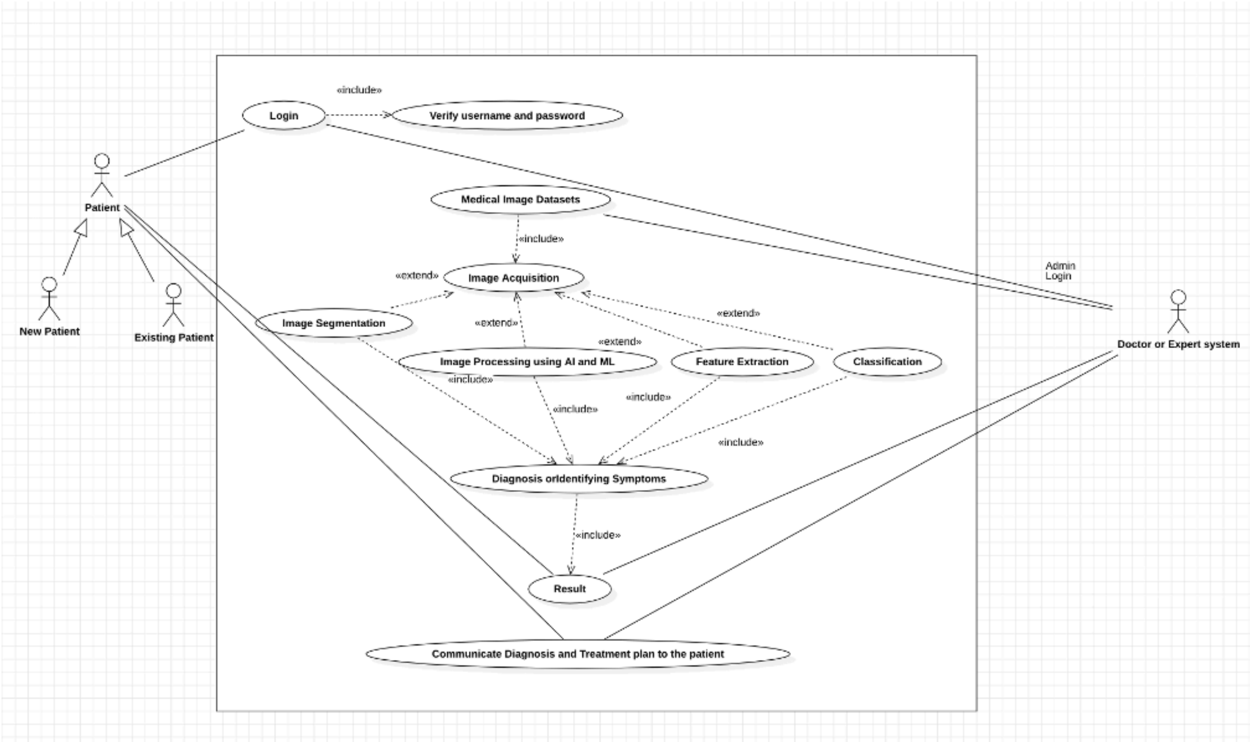

Figure 4 : Use Case Diagram

4.2.3 CLASS DIAGRAM

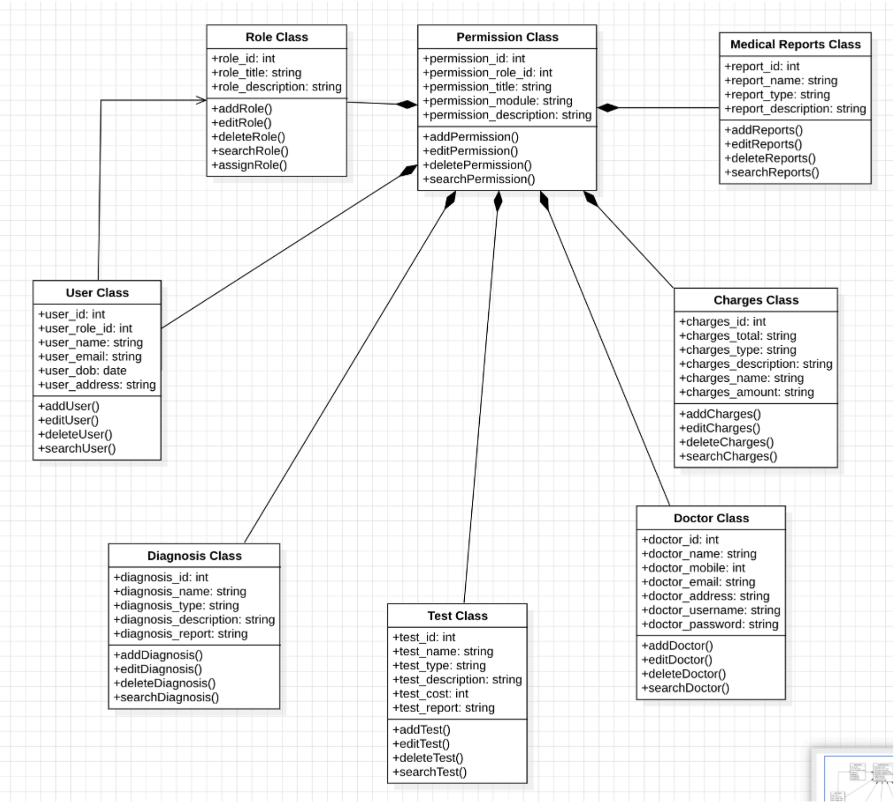

Figure 5 : Class Diagram

4.2.4 SEQUENCE DIAGRAM

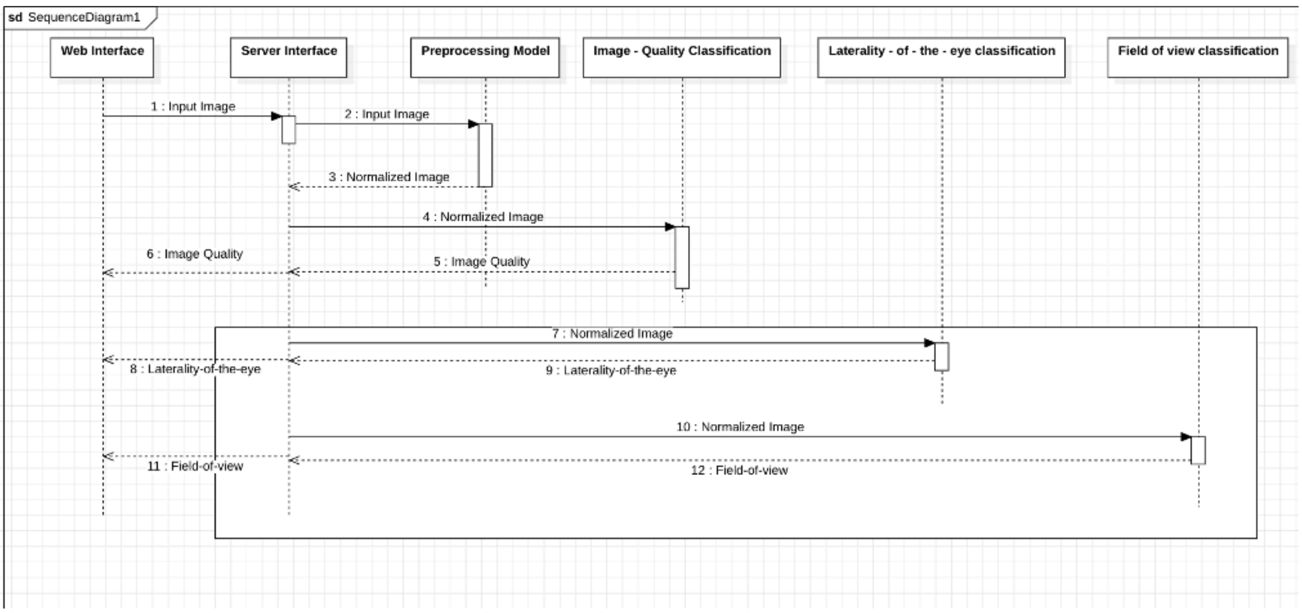

Figure 6 : Sequence Diagram

4.3 CONSTRAINTS, ALTERNATIVES AND TRADEOFFS

Constraints:

1. **Data Privacy and Security**: Medical images contain sensitive patient information, requiring stringent data protection measures to comply with regulations such as HIPAA (in the U.S.) or GDPR (in Europe).
2. **Data Quality and Availability**: High-quality, annotated medical images are crucial for training robust image recognition models. However, the availability of such data can be limited due to privacy concerns, proprietary issues, and variability in medical imaging techniques across facilities.
3. **Computational Resources**: Deep learning models, especially those used in image recognition, require significant computational power and storage, which can be costly and may limit scalability, particularly in resource-limited settings.
4. **Integration with Clinical Workflows**: The technology must integrate seamlessly with existing healthcare IT systems and clinical workflows without causing disruptions or requiring extensive training for medical staff.

Alternatives:

1. **Traditional Machine Learning Models**: Before the widespread adoption of deep learning, traditional machine learning models like support vector machines or decision trees were used. These models generally require less computational power and can be effective for simpler image analysis tasks.
2. **Hybrid Models**: Combining traditional machine learning techniques with deep learning approaches to create hybrid models that leverage the strengths of both can sometimes offer a good balance of performance and computational efficiency.
3. **Transfer Learning**: Using pre-trained models on general image datasets and fine-tuning them for specific medical applications can reduce the need for large domain-specific datasets and lower computational costs.

Trade-offs:

1. **Accuracy vs. Explainability**: Deep learning models, particularly CNNs, often achieve high accuracy in image recognition tasks but are typically considered "black boxes", meaning their decision-making processes are not easily interpretable. This is a significant issue in medical diagnosis, where understanding the rationale behind a diagnosis is as important as the diagnosis itself.
2. **Performance vs. Cost**: Higher model accuracy often requires more complex models with more parameters, which in turn require more data and computational power. This increases costs, making it a significant trade-off for healthcare providers, especially those in low-resource settings.
3. **Data Privacy vs. Model Effectiveness**: Ensuring patient data privacy often means using less data for training models, which can compromise the effectiveness of the models. Techniques like differential privacy or federated learning can be used to address this, but they may introduce complexity or reduce model performance.

5. SCHEDULE, TASKS AND MILESTONES

5.1 GANTT CHART

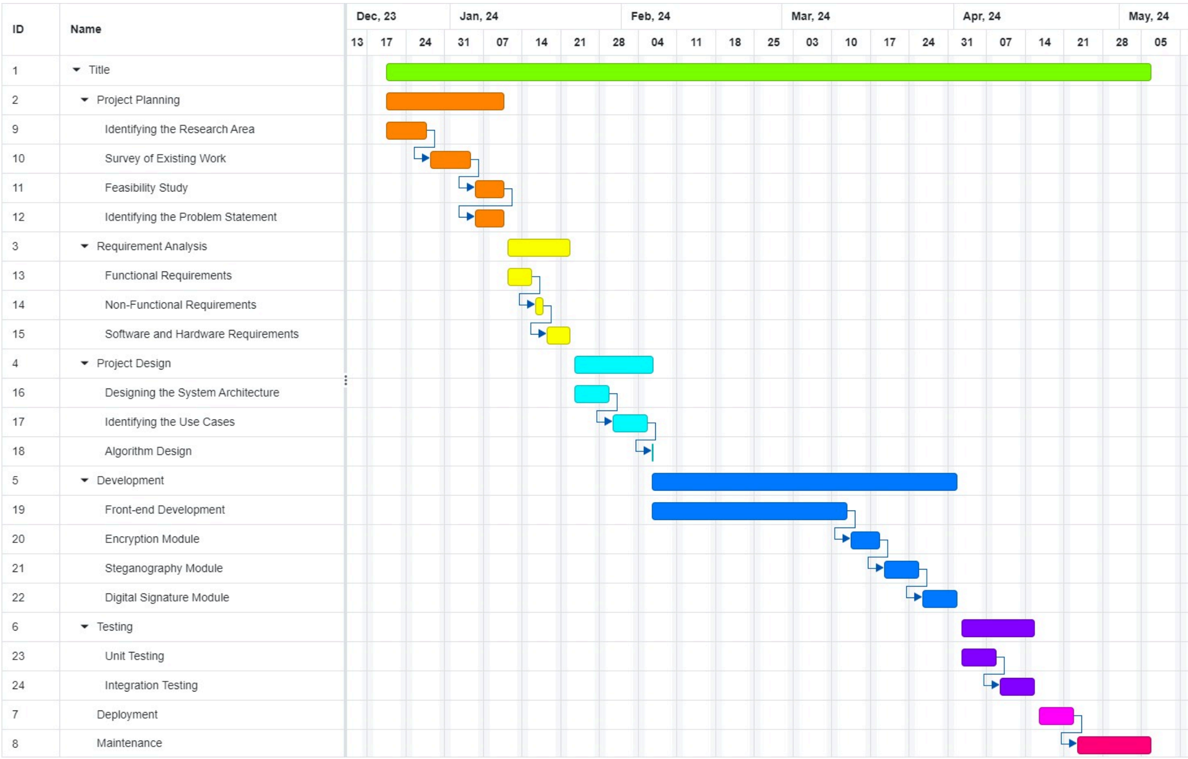

5.2 MODULE DESCRIPTION

5.2.1 Data Ingestion Module

It acts as the entry point for medical images, equipped to handle various formats and preprocess images (normalization, contrast adjustment, anonymization).

5.2.2 Data Storage and Management

It securely stores raw and processed images and metadata, using both cloud and on-premises solutions for optimal accessibility and compliance. Offers a scalable and flexible solution that ensures data integrity and confidentiality.

5.2.3 Image Processing and Feature Extraction

It Enhances image quality and extracts key features for diagnosis, utilizing edge detection and contrast enhancement. Boosts the AI model's accuracy in identifying and diagnosing conditions from images.

5.2.4 AI and Machine Learning Engine

It Employs convolutional neural networks (CNNs) to analyze images and detect disease patterns. Automates diagnostic processes, potentially surpassing human expert accuracy and enhancing diagnostic capabilities.

5.2.5 Decision Support and Reporting Module

It Converts AI analyzes into actionable insights and diagnostic suggestions, complete with reports, confidence scores, and annotated images. Bridges the gap between technical analysis and clinical application, aiding professionals in decision-making.

6. PROJECT DEMONSTRATION

Image Recognition for Medical Diagnosis project would begin with a succinct overview highlighting the system's capabilities, including its ability to process various medical images like X-rays and CT scans using a CNN model to detect medical conditions. The demonstration involves uploading a curated set of medical images to the system, which are then processed in real-time to identify pathological features, with results displayed including condition identifications, confidence scores, and visual markers on the anomalies detected. These results are often accompanied by a rationale, leveraging advanced model features like attention maps for explanation. The demonstration also showcases how the system integrates seamlessly with existing healthcare IT infrastructures, such as Electronic Health Records, and emphasizes robust data security measures compliant with medical data protection standards. A comparison with expert analyses highlights the system's accuracy and efficiency, concluding with a Q&A session to address any queries, underscoring the system’s potential to enhance diagnostic accuracy, reduce diagnostic time, and alleviate workload pressures on medical professionals.

7. RESULTS & DISCUSSION

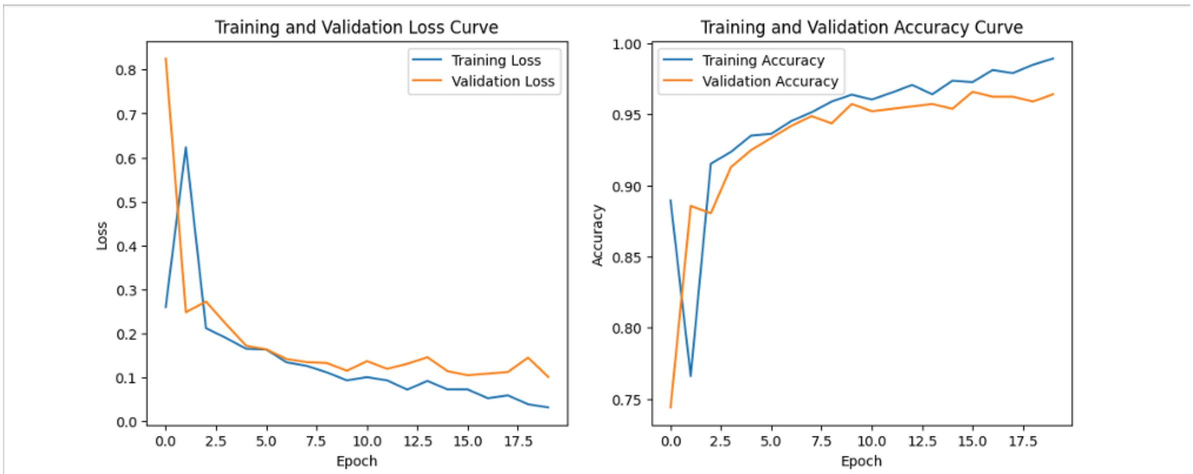

EVALUATE MODEL ON TEST SET

```
#Evaluate Model on Test Set
evaluate_model(model, test_loader, test_indices, 'TEST', criterion, data_path, "ResNet18")
```

```
TEST: Accuracy: 0.9658 | Loss: 0.1004 | Recall: 0.9591 | Precision: 0.9556 | F-score: 0.9573

<Figure size 800x800 with 0 Axes>
```

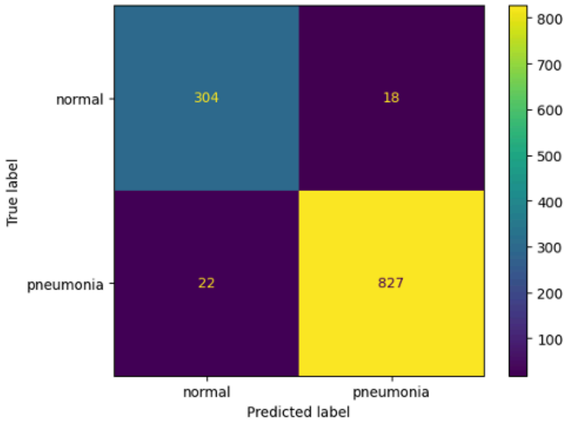


```
#Obtain the TSNE Plot
plot_TSNE(train_loader, device, model)
```

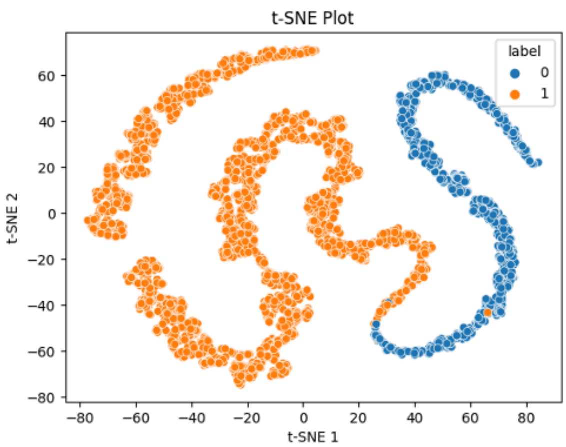


```
#Plot the Within-Class Variance of the dataset
plot_within_class_variance(dataset)
```

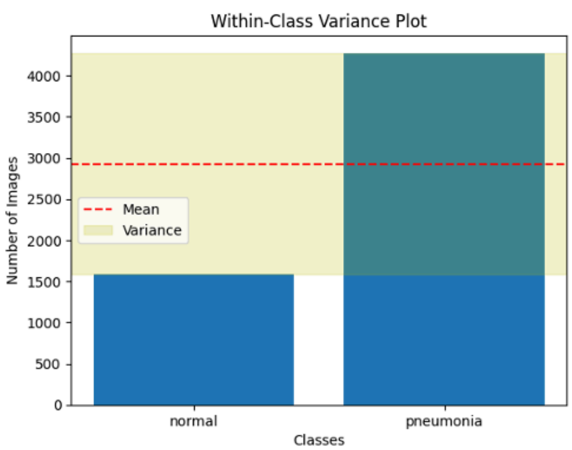

After running the sample code by giving it the datasets of x-ray images, the model detects a 95.23% chance of pneumonia.
This can also be confirmed by the last graph (Within-Class Variance Plot) which shows the mean and variance of a normal x-ray and a one with pneumonia.

8. SUMMARY

Our project showcases the immense potential of image recognition technology in revolutionizing healthcare. Through advanced algorithms and deep learning models, we demonstrate how medical images like X-rays, MRIs, and CT scans can be analyzed accurately to detect diseases and injuries swiftly. This technology empowers clinicians with rapid, precise assessments, paving the way for earlier interventions and personalized treatments. We emphasize the need for ongoing improvements in algorithm accuracy, data privacy, and seamless integration into clinical workflows to maximize benefits for patients and enhance overall healthcare efficacy.

References

[1] Zhang, Q., Zhou, J., & Zhang, B. (Year). Computational Traditional Chinese Medicine diagnosis: A literature survey. Journal Name, Volume (Issue), Page Range. https://doi.org/10.1007/s12065-020-00540-3

[2] Rogers, W., Ryack, B., & Moeller, G. (Year of publication not provided). Computer-aided medical diagnosis: Literature review. Journal of Medical Systems.

[3] Jiang, X., Hu, Z., Wang, S., & Zhang, Y. (2023). Deep Learning for Medical Image-Based Cancer Diagnosis. Cancers, 15(14), 3608. https://doi.org/10.3390/cancers15143608

[4] Li, L.-F., Wang, X., Hu, W.-J., Xiong, N. N., & Du, Y.-X. (2020). Deep Learning in Skin

Disease Image Recognition: A Review. Journal of Medical Systems, Volume 8

[5] Zhang, Q., Zhou, J., & Zhang, B. (Year). Computational Traditional Chinese Medicine diagnosis: A literature survey. Journal Name, Volume (Issue), Page Range. https://doi.org/10.1007/s12065-020-00540-3

[6] P. Gope, Y. Gheraibia, S. Kabir and B. Sikdar, "A secure IoT-based modern healthcare system with fault-tolerant decision-making process", IEEE J. Biomed. Health Informat., vol. 25, no. 3, pp. 862-873, Mar. 2021.

[7] L. Wang, C. Shi, S. Lin, P. Qin and Y. Wang, "Convolutional sparse representation and local density peak clustering for medical image fusion", Int. J. Pattern Recognit. Artif. Intell., vol. 34, no. 7, 2020.

[8] S. Budd, E. C. Robinson and B. Kainz, "A survey on active learning and human-in-the-loop deep learning for medical image analysis", Med. Image Anal., vol. 71, 2021.

[9] W. Chen et al., "Spherical-patches extraction for deep-learning-based critical points detection in 3D neuron microscopy images", IEEE Trans. Med. Imag., vol. 40, no. 2, pp. 527-538, Feb. 2021.

[10] Kwak, D. and Choi, J., 2023. A Method of Breast Cancer Image Diagnosis Using Artificial Intelligence Medical Images Recognition Technology Network. The Journal of Korean Institute of Communications and Information Sciences, 48(2), pp.216-226.

[11] Rana, M., Bhushan, M. Machine learning and deep learning approach for medical image analysis: diagnosis to detection. Multimed Tools Appl 82, 26731–26769 (2023). https://doi.org/10.1007/s11042-022-14305-w

[12] 1. Z. Guo, Y. Shen, S. Wan, W. -L. Shang and K. Yu, "Hybrid Intelligence-Driven Medical Image Recognition for Remote Patient Diagnosis in Internet of Medical Things," in IEEE Journal of Biomedical and Health Informatics, vol. 26, no. 12, pp. 5817-5828, Dec. 2022, doi: 10.1109/JBHI.2021.3139541.

[13] Citation: Samee, N.A.; Mahmoud, N.F.; Atteia, G.; Abdallah, H.A.; Alabdulhafith, M.; Al-Gaashani, M.S.A.M.; Ahmad, S.; Muthanna, M.S.A. Classification Framework for Medical Diagnosis of Brain Tumor with an Effective Hybrid Transfer Learning Model. Diagnostics 2022, 12, 2541. https://doi.org/10.3390/diagnostics12102541

APPENDIX A – SAMPLE CODE

This code is written in Python using ResNet18 on Jupyter Notebook

Datasets used:

https://www.kaggle.com/datasets/paulti/chest-xray-images

https://dataverse.harvard.edu/file.xhtml?fileId=5194823&version=5.1

https://www.kaggle.com/datasets/nihchest-xrays/data

https://drive.google.com/drive/folders/10e-Yf_PxUTCDvh97mU1C5dp5tz7flEKk?usp=sharing

```
[ ]: from google.colab import drive
     drive.mount('/content/drive', force_remount = True)

     Mounted at /content/drive

[ ]: import os, time, random, torch, warnings
     import numpy as np
     from PIL import Image
     import torch.nn as nn
     from tqdm import tqdm
     import torch.optim as optim
     import matplotlib.pyplot as plt
     import torchvision.datasets as datasets
     import torchvision.transforms as transforms
     from torch.utils.data import DataLoader, random_split
     from sklearn.metrics import precision_score, recall_score, f1_score

     warnings.simplefilter("ignore")
     device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
     device

[2]: device(type='cuda')

[ ]: data_path = "/content/drive/My Drive/Simple_Chest_XRay/"
     #data_path = "/content/drive/My Drive/NIH_Chest_XRay/"
     #data_path = "/content/drive/My Drive/Harvard_Chest_XRay/"

     sample_ratio = 1
     batch_size = 96
     num_epochs = 20

[ ]: # Set random seed for reproducibility
     torch.manual_seed(42)
     np.random.seed(42)
     random.seed(42)
```

```
[ ]: def save_metrics(loss, accuracy, model):
       np.save("{}{}_train_loss.npy".format(data_path, model), loss)
       np.save("{}{}_train_accuracy.npy".format(data_path, model), accuracy)
```

DATA PREPROCESSING

```
[ ]: %run "/content/drive/My Drive/Colab Notebooks/utils.ipynb"
     dataset, train_loader, train_indices, test_loader, test_indices, val_loader, val_indices = data_preprocess(data_path, sample_ratio, batch_size)
```

DOWNLOAD RESNET18 MODEL AND TRAIN

```
[ ]: # Define the ResNet18 model and set Pretraining to False to train model from scratch
     model = torch.hub.load('pytorch/vision:v0.9.0', 'resnet18', pretrained = False)
     model.fc = nn.Linear(512, len(dataset.classes))
     model.to(device)

     # Define loss function a  s CrossEntropy and optimizer as Adam Optimizer
     criterion = nn.CrossEntropyLoss()
     optimizer = optim.Adam(model.parameters(), lr = 0.0001)

     losses, accuracies, v_accuracies, v_losses = train_model(model, criterion, optimizer, "ResNet18", num_epochs)
```

```
Using cache found in /root/.cache/torch/hub/pytorch_vision_v0.9.0
Epoch 1/20: 100%|██████████| 43/43 [01:21<00:00,  1.89s/batch, Accuracy=0.889, Loss=0.26, Precision=0.852, Recall=0.883, F1 Score=0.865]
VALIDATION: Accuracy: 0.7440 | Loss: 0.8251 | Recall: 0.5482 | Precision: 0.8684 | F-score: 0.5122

Epoch 2/20: 100%|██████████| 43/43 [01:18<00:00,  1.83s/batch, Accuracy=0.766, Loss=0.623, Precision=0.781, Recall=0.774, F1 Score=0.778]
VALIDATION: Accuracy: 0.8857 | Loss: 0.2483 | Recall: 0.8255 | Precision: 0.8850 | F-score: 0.8482

Epoch 3/20: 100%|██████████| 43/43 [01:19<00:00,  1.85s/batch, Accuracy=0.915, Loss=0.212, Precision=0.818, Recall=0.814, F1 Score=0.816]
VALIDATION: Accuracy: 0.8805 | Loss: 0.2726 | Recall: 0.7928 | Precision: 0.9200 | F-score: 0.8298

Epoch 4/20: 100%|██████████| 43/43 [01:19<00:00,  1.85s/batch, Accuracy=0.924, Loss=0.19, Precision=0.84, Recall=0.834, F1 Score=0.837]
VALIDATION: Accuracy: 0.9130 | Loss: 0.2211 | Recall: 0.9265 | Precision: 0.8825 | F-score: 0.8991

Epoch 5/20: 100%|██████████| 43/43 [01:20<00:00,  1.87s/batch, Accuracy=0.935, Loss=0.165, Precision=0.856, Recall=0.85, F1 Score=0.853]
VALIDATION: Accuracy: 0.9249 | Loss: 0.1720 | Recall: 0.8839 | Precision: 0.9286 | F-score: 0.9027

Epoch 6/20: 100%|██████████| 43/43 [01:22<00:00,  1.92s/batch, Accuracy=0.936, Loss=0.164, Precision=0.867, Recall=0.861, F1 Score=0.864]
VALIDATION: Accuracy: 0.9334 | Loss: 0.1638 | Recall: 0.8953 | Precision: 0.9392 | F-score: 0.9139

Epoch 7/20: 100%|██████████| 43/43 [01:21<00:00,  1.90s/batch, Accuracy=0.945, Loss=0.135, Precision=0.876, Recall=0.87, F1 Score=0.873]
VALIDATION: Accuracy: 0.9420 | Loss: 0.1420 | Recall: 0.9158 | Precision: 0.9395 | F-score: 0.9266

Epoch 8/20: 100%|██████████| 43/43 [01:19<00:00,  1.85s/batch, Accuracy=0.951, Loss=0.126, Precision=0.884, Recall=0.879, F1 Score=0.881]
VALIDATION: Accuracy: 0.9488 | Loss: 0.1348 | Recall: 0.9351 | Precision: 0.9384 | F-score: 0.9367

Epoch 9/20: 100%|██████████| 43/43 [01:22<00:00,  1.93s/batch, Accuracy=0.959, Loss=0.112, Precision=0.892, Recall=0.886, F1 Score=0.889]
VALIDATION: Accuracy: 0.9437 | Loss: 0.1329 | Recall: 0.9243 | Precision: 0.9357 | F-score: 0.9298

Epoch 10/20: 100%|██████████| 43/43 [01:22<00:00,  1.91s/batch, Accuracy=0.964, Loss=0.0935, Precision=0.898, Recall=0.893, F1 Score=0.895]
VALIDATION: Accuracy: 0.9573 | Loss: 0.1153 | Recall: 0.9466 | Precision: 0.9482 | F-score: 0.9474

Epoch 11/20: 100%|██████████| 43/43 [01:21<00:00,  1.89s/batch, Accuracy=0.96, Loss=0.101, Precision=0.903, Recall=0.898, F1 Score=0.9]
VALIDATION: Accuracy: 0.9522 | Loss: 0.1371 | Recall: 0.9302 | Precision: 0.9509 | F-score: 0.9398

Epoch 12/20: 100%|██████████| 43/43 [01:23<00:00,  1.94s/batch, Accuracy=0.965, Loss=0.0934, Precision=0.907, Recall=0.903, F1 Score=0.905]
VALIDATION: Accuracy: 0.9539 | Loss: 0.1196 | Recall: 0.9314 | Precision: 0.9540 | F-score: 0.9419

Epoch 13/20: 100%|██████████| 43/43 [01:22<00:00,  1.91s/batch, Accuracy=0.971, Loss=0.0724, Precision=0.912, Recall=0.907, F1 Score=0.909]
VALIDATION: Accuracy: 0.9556 | Loss: 0.1312 | Recall: 0.9454 | Precision: 0.9454 | F-score: 0.9454

Epoch 14/20: 100%|██████████| 43/43 [01:22<00:00,  1.92s/batch, Accuracy=0.964, Loss=0.0921, Precision=0.915, Recall=0.91, F1 Score=0.913]
VALIDATION: Accuracy: 0.9573 | Loss: 0.1460 | Recall: 0.9338 | Precision: 0.9605 | F-score: 0.9459

Epoch 15/20: 100%|██████████| 43/43 [01:25<00:00,  1.98s/batch, Accuracy=0.974, Loss=0.0728, Precision=0.918, Recall=0.914, F1 Score=0.916]
VALIDATION: Accuracy: 0.9539 | Loss: 0.1144 | Recall: 0.9515 | Precision: 0.9376 | F-score: 0.9442
```

```
Epoch 16/20: 100%|██████████| 43/43 [01:23<00:00,  1.94s/batch, Accuracy=0.973, Loss=0.0728, Precision=0.921, Recall=0.917, F1 Score=0.919]
VALIDATION: Accuracy: 0.9659 | Loss: 0.1051 | Recall: 0.9489 | Precision: 0.9665 | F-score: 0.9572

Epoch 17/20: 100%|██████████| 43/43 [01:24<00:00,  1.96s/batch, Accuracy=0.981, Loss=0.0529, Precision=0.925, Recall=0.92, F1 Score=0.923]
VALIDATION: Accuracy: 0.9625 | Loss: 0.1088 | Recall: 0.9611 | Precision: 0.9484 | F-score: 0.9544

Epoch 18/20: 100%|██████████| 43/43 [01:24<00:00,  1.96s/batch, Accuracy=0.979, Loss=0.0593, Precision=0.928, Recall=0.923, F1 Score=0.925]
VALIDATION: Accuracy: 0.9625 | Loss: 0.1126 | Recall: 0.9392 | Precision: 0.9682 | F-score: 0.9523

Epoch 19/20: 100%|██████████| 43/43 [01:25<00:00,  1.99s/batch, Accuracy=0.985, Loss=0.0392, Precision=0.93, Recall=0.926, F1 Score=0.928]
VALIDATION: Accuracy: 0.9590 | Loss: 0.1452 | Recall: 0.9350 | Precision: 0.9637 | F-score: 0.9480

Epoch 20/20: 100%|██████████| 43/43 [01:26<00:00,  2.00s/batch, Accuracy=0.989, Loss=0.0319, Precision=0.933, Recall=0.929, F1 Score=0.931]
VALIDATION: Accuracy: 0.9642 | Loss: 0.1013 | Recall: 0.9604 | Precision: 0.9523 | F-score: 0.9563

[0.26026202372208723, 0.6233695970603274, 0.2120919976402533, 0.18972522627801305, 0.16519693304312233, 0.16359962942499043, 0.1345113130873894, 0.126482
46012524472, 0.11169907397236815, 0.0934581851809817, 0.10083207577488777, 0.09339687339807028, 0.07239830867637509, 0.09214521172563458, 0.0728305868520
9001, 0.07281481958952989, 0.05288402761634195, 0.059270327886696936, 0.03919073106472479, 0.03190121432472312] [0.8251181897856676, 0.24834805340490243,
0.27261529450937344, 0.22111985748538385, 0.17201383701770379, 0.16376029956869706, 0.14200975234801452, 0.13477527516108934, 0.13291411639633033, 0.1153
3337821325751, 0.13713266377562955, 0.11961223058078324, 0.13117166206641154, 0.14599780996454048, 0.11439624133777293, 0.1050591233342174, 0.10877843603
885927, 0.11255094560289139, 0.14516717344109262, 0.1013102003535313]
```

SAVE MODEL PARAMETERS

```
[ ]: torch.save(model.state_dict(), "{}resnet18.pth".format(data_path))
```

```
[ ]: #Plot the Model Accuracy and Loss on Training and Validation Sets
plot_model_curves(losses, accuracies, v_accuracies, v_losses)
```

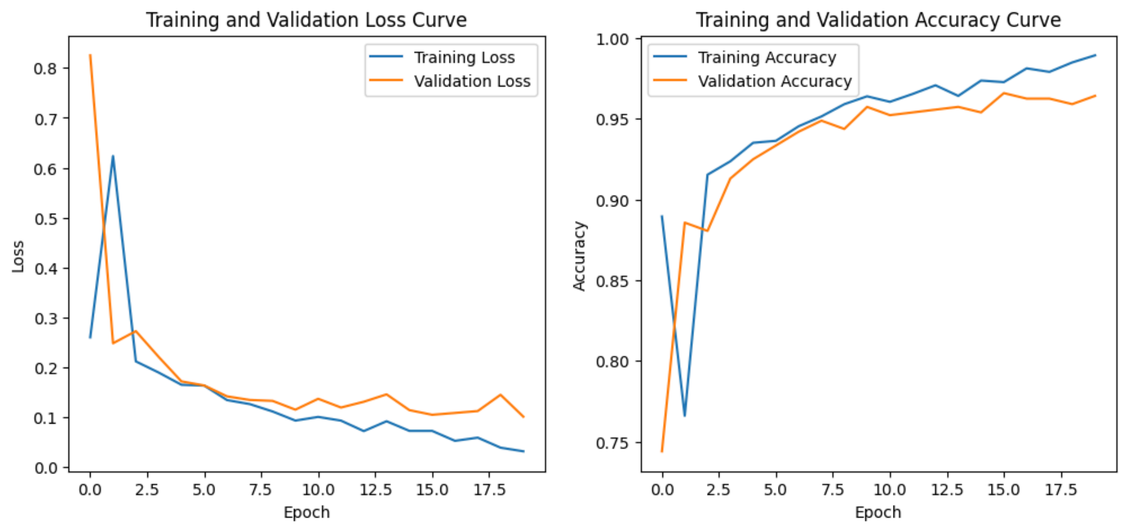

EVALUATE MODEL ON TEST SET

```
[ ]: #Evaluate Model on Test Set
evaluate_model(model, test_loader, test_indices, 'TEST', criterion, data_path, "ResNet18")
```

```
TEST: Accuracy: 0.9658 | Loss: 0.1004 | Recall: 0.9591 | Precision: 0.9556 | F-score: 0.9573

<Figure size 800x800 with 0 Axes>
```

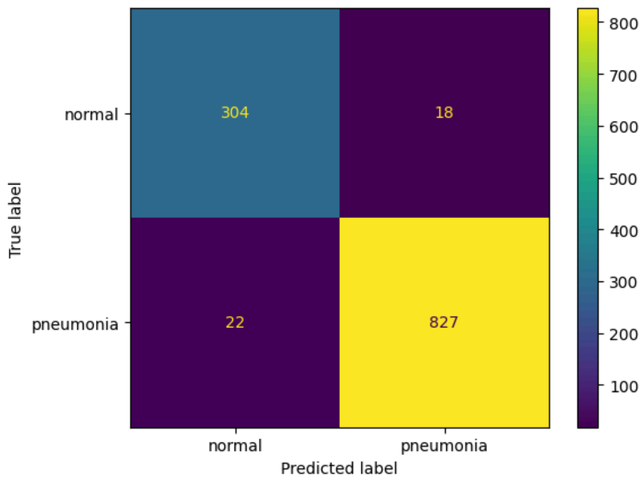

```
[ ]: #Obtain the TSNE Plot
     plot_TSNE(train_loader, device, model)
```

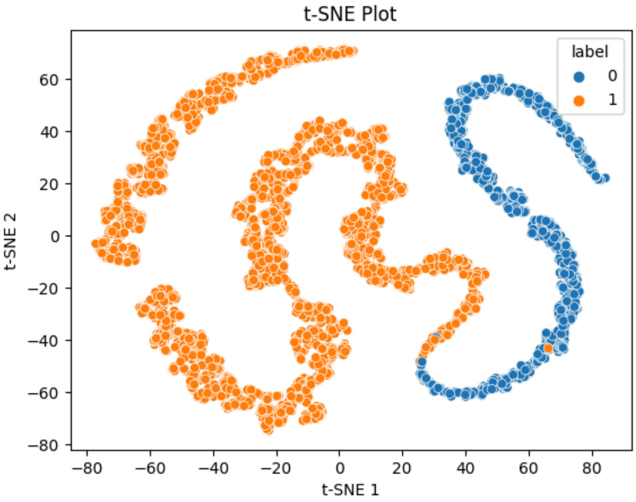


```
[ ]: #Plot the Within-Class Variance of the dataset
     plot_within_class_variance(dataset)
```